

Paraphrase Detection

Jens Christian Valen Leynse

Computer Science / 2024

jcleynse@stud.ntnu.no

and

Vincenzo Spinello

Computer Science / 2024

vincensp@stud.ntnu.no

Abstract

In recent years, significant advancements have emerged in the field of Artificial Intelligence, particularly in the domains of deep learning and Natural Language Processing (NLP). This study delves into the application of these cutting-edge techniques to the task of paraphrase detection, a fundamental challenge in NLP. Specifically, the research explores the fine-tuning of pre-trained language models such as ALBERT and compares their performance against traditional machine learning approaches like Support Vector Machines (SVMs). By investigating the efficacy of these methodologies, this study contributes to the understanding of their suitability for addressing real-world NLP tasks, offering insights into their strengths and limitations in paraphrase detection. Despite efforts to optimize both approaches, achieving a validation accuracy of around 72% poses a challenge. The study concludes with a reflection on the implications of these findings for the task of paraphrase detection and highlights the need for further investigation into alternative methods to enhance model performance.

1 Introduction

Paraphrase Detection (PD) is the task of establishing if two sentences convey the same meaning, with significant implications for various natural language processing (NLP) applications. In the Information Retrieval (IR) field, accurately identifying paraphrased sentences can lead to improved search results. PD models also play a crucial role in evaluating paraphrase generation models.

Another important application of PD is in Text Summarization where detecting redundant sentences aids in producing concise and meaningful summary. However, Paraphrase detection presents a set of challenges, requiring models capable of deep, context-aware analysis to accurately disambiguate meanings. Additionally, the quality of the training dataset is essential for developing high-performance models, as noisy data can impact model effectiveness.

In this work, we propose a Transformer-based model, specifically the ALBERT (A Lite BERT for Self-supervised Learning of Language Representations) model¹, developed by Lan et al. (2020). Transformers, as described in Turner (2024) and Vaswani et al. (2023), have emerged as state-of-the-art for many NLP tasks due to their self-attention mechanism, which enables semantic and contextual understanding.

Furthermore, we employ Transfer Learning, as described by Zhuang et al. (2020) and Iman et al. (2022), leveraging a pre-trained generic model and fine-tuning it for the specific domain. In our approach, we freeze most layers of ALBERT and fine-tune only the last ones to gain domain-specific knowledge while retaining generic knowledge from pre-training. Compared to BERT, ALBERT offers a more efficient fine-tuning process due to its smaller size and parameter count.

A Support Vector Machine model will also be trained from scratch and its results will be compared against the fine-tuned model.

¹ALBERT is available at this link: https://huggingface.co/docs/transformers/en/model_doc/albert

2 Background

In this section, we provide the necessary background information to understand the theoretical basis and contributions of our project. Our aim is to introduce terminology and key concepts that will be used throughout the document, ensuring clarity and coherence for readers who may not be familiar with the topic.

- **Word Embeddings and Positional Encoding:** Word embeddings, as discussed in previous research Mikolov et al. (2013), are foundational concepts in natural language processing, crucial for representing the semantic content of tokens. These embeddings capture the contextual meaning of words in a continuous vector space, enabling computational models to process textual data effectively. Additionally, we discuss the importance of positional encoding, which enriches token embeddings with information about their positions in the input sequence, ensuring that transformer-based models can effectively process sequential data while preserving positional relationships.
- **Transformer Architecture:** The transformer architecture represents a state-of-the-art model for various natural language processing tasks. Transformers leverage self-attention mechanisms and multi-layer neural networks to process input sequences and capture complex patterns and dependencies.
- **Incorporating Positional Information:** Incorporating positional information into token embeddings is a crucial step in transformer-based models. This process ensures that the model maintains awareness of token positions within the input sequence, enabling effective processing of sequential data. By discussing the methods for incorporating positional information, we lay the groundwork for understanding the mechanisms underlying transformer-based models and their applications.
- **Self Attention:** Also known as intra-attention, this mechanism aims to establish relationships between different positions within a sequence in order to compute a representation of the entire sequence.
- **Support Vector Machine (SVM):** SVM is a supervised learning algorithm that can be used for classification or regression tasks. Unlike Transformer-based architectures, which rely on self-attention mechanisms and neural networks, SVM operates by finding the optimal hyperplane that separates data points into different classes based on their features. Using kernels, data points are projected into higher-dimensional feature spaces, enabling linear separability. In this work, an SVM is used as a baseline to compare the performance of the ALBERT model.

2.1 Dataset

The dataset comprises 400,000 pairs of questions sourced from Quora, a popular question-and-answer platform. Each pair consists of two questions (question1 and question2) along with their corresponding IDs, indicating the uniqueness of each pair. Additionally, the dataset includes labels denoted by the "is_duplicate" column, which indicates whether the questions in each pair are semantically similar or not (0 for dissimilar, 1 for similar). This diverse collection of question pairs covers a wide range of topics and language patterns, making it a valuable resource for studying natural language understanding and semantic similarity. Through analysis of the syntactic and distributional characteristics of this dataset, we can gain deeper insights into the linguistic nuances of human-generated text and develop robust models for various NLP tasks, such as duplicate question detection and semantic similarity assessment.




Figure 1: >50000

The provided output offers a comprehensive view of the linguistic and distributional characteristics of the Quora sentence pairs dataset. The "Distribution of Part-of-Speech Tags" figure reveals insights into the syntactic composition of the text data. Notably, nouns (NN) and plural nouns (NNS) emerge as the most frequent part-of-speech tags, indicating a prevalence of subject matter and potentially informative phrases within the dataset. Additionally, the high occurrence of adjectives (JJ) and adverbs (RB) suggests a descriptive nature to the sentences, hinting at rich contextual information that may influence semantic similarity.




Figure 2:

The "Histogram values" figure illustrates the distribution of similarity scores between sentence pairs using a histogram. This visualization reveals a diverse spectrum of similarity scores, emphasizing intervals where scores are most concentrated. As scores deviate from these peaks, their frequency gradually diminishes, reflecting the dataset's semantic diversity. This insight underscores the varying degrees of relatedness among sentence pairs.




Figure 3: Data Characteristics

The "Density values" section displays a smoothed distribution of similarity scores via kernel density estimation, offering insights into the dataset's characteristics. Peaks in density at extremes indicate a higher proportion of very similar or dissimilar sentence pairs.

3 Related Work

Paraphrase identification has been extensively explored in natural language processing, with researchers investigating various methodologies to tackle this challenge. In recent years, transformer-based models have emerged as powerful tools for capturing semantic relationships between sentences.

Khairova et al. (2022) conducted a study on paraphrase identification using the BERT model and its variants. They specifically explored the effectiveness of the Sentence-BERT model, which is fine-tuned for sentence embeddings, in classifying paraphrased sentences within news corpora. Their work demonstrated the advantages of utilizing fine-tuned BERT models for semantic similarity tasks, showcasing promising results in identifying paraphrased sentences from news articles related to COVID-19.

In contrast, our work focuses on leveraging the ALBERT model for paraphrase detection tasks. ALBERT, an evolution of the BERT model, introduces innovations aimed at enhancing efficiency and scalability while maintaining high performance. We aim to explore the applicability of ALBERT in the context of paraphrase detection and compare its performance against other approaches.

Khairova et al. (2022) employed the Sentence-BERT model for paraphrase identification, we propose fine-tuning the ALBERT model for domain-specific paraphrase detection tasks. This approach allows us to adapt the pre-trained ALBERT model to our specific domain, potentially improving the accuracy and robustness of paraphrase detection in our target application.

Additionally, recent research in natural language understanding has delved into methodologies for common sense verification. Notably, Li et al. (2020) investigated discerning meaningful from meaningless natural language statements using the lightweight ALBERT model based on the Transformer mechanism. While their task primarily revolves around distinguishing between common sense and extraordinary knowledge, the methodologies employed hold relevance to paraphrase detection. Both tasks necessitate deep semantic understanding and contextual analysis of language, positioning the ALBERT model as a promising candidate for exploration in our paraphrase detection endeavor.




Figure 4: The Transformer block (adapted from Turner (2024)).

4 Architecture

In this section, we provide an overview of the ALBERT model architecture, preceded by a brief introduction to the general Transformer architecture.

4.1 Transformer Model

The Transformer architecture represents a significant breakthrough in natural language processing (NLP), introducing a novel approach that relies on self-attention mechanisms to capture contextual dependencies within input sequences. At its core, the Transformer model comprises two fundamental components: the Self-Attention mechanism and the Feed-Forward Neural Network. The Self-Attention mechanism allows the model to weigh the importance of different words in a sequence based on their contextual relevance. By considering all tokens simultaneously, Self-Attention enables the model to capture long-range dependencies and contextual nuances more effectively than traditional recurrent neural networks (RNNs) or convolutional neural networks (CNNs). In addition to the Self-Attention mechanism, the Transformer architecture incorporates a Feed-Forward Neural Network, which processes the information extracted by the Self-Attention mechanism to generate meaningful representations. This network consists of multiple layers of fully connected feed-forward units, each equipped with activation functions to introduce non-linearity and facilitate complex feature extraction. To enhance training stability and facilitate gradient flow, the Transformer architecture employs residual connections, discussed in He et al. (2015), around each sub-layer, including the Self-Attention mechanism and the Feed-Forward Neural Network. These residual connections ensure that the original input to a layer is preserved and added back to the transformed output of that layer. This mechanism helps in addressing the vanishing gradient problem by ensuring that the gradient can flow directly through the identity connection during backpropagation, thereby facilitating the training of deep neural networks. Furthermore, layer normalization is applied after each sub-layer to ensure consistent training dynamics and improve overall model performance. The Transformer block is depicted in figure 4 which illustrates the core components and connections within the architecture




Figure 5: Architecture of the ALBERT model (adapted from Li et al. (2020)).

4.2 ALBERT Model

With a foundational understanding of Transformer architecture established, we now turn our attention to the specifics of the ALBERT model. Developed as an evolution of the BERT (Bidirectional Encoder Representations from Transformers) model, ALBERT introduces several innovations aimed at enhancing efficiency and scalability while maintaining state-of-the-art performance in various NLP tasks. ALBERT improves upon BERT by reducing its parameter count while maintaining or even enhancing its performance across various natural language processing tasks. The ALBERT architecture is depicted in figure 5. The following are the key components of ALBERT:

- **Embeddings Layer:** ALBERT starts with an embedding layer that converts input tokens into dense vector representations. These embeddings capture the semantic content of tokens and their positional relationships within the input sequence.
- **Transformer Encoder:** The core of the ALBERT model is a multi-layer bidirectional Transformer encoder. Each encoder layer consists of self-attention mechanisms and feedforward neural networks, allowing the model to capture complex patterns and dependencies within the input sequence.
- **Parameter Reduction Techniques:** ALBERT introduces several parameter reduction techniques to achieve efficiency without sacrificing performance. These techniques include factorized embedding parameterization, cross-layer parameter sharing, and inter-sentence coherence loss.

ALBERT is pre-trained using unsupervised learning objectives, such as masked language modeling (MLM) and next sentence prediction (NSP). The MLM objective involves randomly masking tokens in the input sequence and training the model to predict the masked tokens based on their context. The NSP objective focuses on predicting whether two consecutive sentences in a text are contiguous or randomly sampled from the corpus.

5 Experiments and Results

In this section, we will discuss the experiments conducted to address the Paraphrase Detection task. We will outline the experimental setup, including the models used, training procedures, and evaluation metrics. Additionally, the results and analysis of each experiment will be presented to provide insights into the performance of the models.

5.1 Experimental Setup

To conduct the experiments, we fine-tuned the ALBERT-base-v2 model using the 'AutoModel' package. We trained two variations of the model:

- **ALBERT-base-v2 (200k samples):** This model was trained for 5 epochs using 200,000 samples, approximately 50% of the total dataset. We selected a learning rate of 0.0005, a batch size of 32, and a maximum input length of 128 tokens for this model.
- **ALBERT-base-v2 (400k samples):** For this model, we utilized the entire dataset consisting of 400,000 samples and trained it for 2 epochs. Similar to the previous model, we maintained a learning rate of 0.0005, a batch size of 32, and a maximum input length of 128 tokens.

In our experiments, we employed Adam as the optimizer and BCEWithLogitsLoss as the loss function. Adam optimizer is widely recognized for its effectiveness in training deep learning models due to its adaptive learning rate method. Similarly, BCEWithLogitsLoss, which combines a sigmoid activation function and binary cross-entropy loss, is commonly used for binary classification tasks like paraphrase detection. Together, these choices allowed us to effectively train our models and optimize them for the specific task at hand

5.2 Experimental Results

- **Support Vector Machine (SVM) Results:**

Validation Accuracy: 0.7497

Validation Precision: 0.7053

Validation Recall: 0.5306

Validation F1 Score: 0.6056

The SVM model achieved a respectable validation accuracy of 74.97%. It demonstrated balanced precision and recall scores, with a precision of 70.53% and a recall of 53.06%. The F1 score, which considers both precision and recall, was calculated to be 60.56%.

Model	validation		training		test	
	loss	accuracy	loss	accuracy	loss	accuracy
ALBERT-base-v2 (200k samples)	54.13%	70.29%	57.85%	67.63%	52.78%	71.84%
	52.65%	72.07%	54.55%	70.29%		
	52.24%	72.17%	53.71%	71.31%		
	52.29%	72.55%	53.16%	71.66%		
	52.45%	72.17%	52.74%	72.06%		
ALBERT-base-v2 (400k samples)	52.23%	71.14%	56.28%	67.00%	51.36%	71.63%
	51.34%	71.73%	53.53%	69.19%		

Table 1: Model Performance Metrics

- **ALBERT-base-v2 (200k samples) Results:**

Training Loss: Decreased from 0.5785 to 0.5274 over 5 epochs.

Training Accuracy: Increased from 0.6763 to 0.7206 over 5 epochs.

Validation Loss: Decreased from 0.5413 to 0.5245 over 5 epochs.

Validation Accuracy: Increased from 0.7029 to 0.7217 over 5 epochs.

The ALBERT-base-v2 (200k samples) model showcased promising performance during training and validation. Over the course of 5 epochs, the model demonstrated a consistent decrease in training and validation loss, accompanied by an increase in accuracy. The final test accuracy reached 71.84%, indicating robust generalization to unseen data.




Figure 6: Performance Evaluation during Training and Validation

- **ALBERT-base-v2 (400k samples) Results:**

Training Loss: Decreased from 0.5628 to 0.5353 over 2 epochs.

Training Accuracy: Increased from 0.6700 to 0.6919 over 2 epochs.

Validation Loss: Decreased from 0.5223 to 0.5134 over 2 epochs.

Validation Accuracy: Increased from 0.7114 to 0.7173 over 2 epochs.

The ALBERT-base-v2 (400k samples) exhibited similar trends in training and validation as the base model, albeit over a shorter training period of 2 epochs. Despite the increased model complexity and longer training time, the performance improvement was marginal, with a test accuracy of 71.63%. Further optimization and fine-tuning may be required to fully leverage the benefits of the larger model architecture.

6 Evaluation and Discussion

The SVM model exhibited the highest validation accuracy among all models, indicating its effectiveness. This highlights the robustness of traditional machine learning approaches like SVM in handling text classification tasks. However, despite its high validation accuracy, the SVM model may lack the ability to capture complex semantic relationships present in natural language data.

In contrast, the ALBERT classifiers, both the base and large models, demonstrated notable improvements in accuracy over epochs during training. Despite starting with lower validation accuracies compared to the SVM model, the ALBERT models showcased consistent performance enhancement throughout the training process. This suggests that the ALBERT models have the capacity to learn intricate patterns and semantic representations embedded within the dataset, leveraging the power of transformer-based architectures for natural language understanding tasks.

Although the final validation accuracy of the ALBERT models was slightly lower than that of the SVM model, the increasing accuracy trends observed during training indicate their potential for further enhancement. With additional fine-tuning and optimization strategies, the ALBERT models may surpass the performance of traditional SVM classifiers. This underscores the importance of leveraging state-of-the-art deep learning techniques for advancing the field of natural language processing and text classification.

6.1 What Went Wrong and Why

Despite the overall promising performance of the ALBERT models, there are several factors that may have contributed to suboptimal outcomes:

1. **Limited Data Augmentation:** The dataset size, although substantial with 400,000 question pairs, may not have been sufficient to fully train the ALBERT models, especially the larger variant. Augmenting the dataset with additional question pairs or employing more sophisticated data augmentation techniques could potentially improve model performance.
2. **Hyperparameter Tuning:** The choice of hyperparameters, such as learning rate and batch size, significantly impacts model training and convergence. Suboptimal hyperparameter settings may have hindered the ALBERT models' ability to reach their full potential. Conducting a more exhaustive hyperparameter search could lead to better performance.
3. **Model Complexity:** The ALBERT large model, despite its enhanced capacity to capture complex patterns, may have suffered from overfitting due to its increased complexity. Regularization techniques such as dropout or weight decay could mitigate overfitting and improve generalization.
4. **Lack of Fine-Tuning:** Pre-trained transformer models like ALBERT often require fine-tuning on task-specific data to adapt their representations effectively. Insufficient fine-tuning or improper task-specific adaptation may have limited the models' performance. Fine-tuning strategies tailored to the dataset and task at hand could yield better results.

Addressing these challenges through data augmentation, hyperparameter tuning, regularization, and fine-tuning strategies can potentially enhance the ALBERT models' performance and bridge the gap between their current performance and the desired outcome. Additionally, leveraging larger datasets and exploring alternative model architectures may offer further avenues for improvement.

7 Conclusion and Future Work

In this study, we investigated the application of state-of-the-art techniques in Artificial Intelligence, particularly in the field of Natural Language Processing (NLP), to the challenging task of paraphrase detection. By fine-tuning pre-trained language models like ALBERT and comparing their performance against traditional machine learning methods such as Support Vector Machines (SVMs), we gained valuable insights into their efficacy and limitations. Our experiments revealed that while both approaches showed promise, achieving a validation accuracy of approximately 72% posed a significant challenge. This underscores the complexity of paraphrase detection and highlights the need for further research to improve model performance. Despite the limitations encountered, our study contributes to the growing body of knowledge in NLP by shedding light on the strengths and weaknesses of different methodologies. Moving forward, we advocate for continued exploration of alternative techniques, including advanced model architectures, novel feature engineering strategies, and more extensive data augmentation methods. Further investigation into the performance of alternative language models, such as BERT, RoBERTa, or XLNet, could provide additional insights into their suitability for paraphrase detection tasks. To conclude, our study highlights the critical need for continued research in NLP and paraphrase detection, emphasizing our efforts to advance the development of resilient and precise models capable of adeptly navigating the intricacies of human language.

References

- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition, 2015. URL <https://arxiv.org/pdf/1512.03385.pdf>.
- Mohammadreza Iman, Khaled Rasheed, and Hamid Reza Arabnia. A review of deep transfer learning and recent advancements, 2022. URL <https://arxiv.org/pdf/2201.09679.pdf>.
- Nina Khairova, Anastasiia Shapovalova, Orken Mamyrbayev, Nataliia Sharonova, and Kuralay Mukhsina. Using bert model to identify sentences paraphrase in the news corpus, 2022. URL <https://ceur-ws.org/Vol-3171/paper6.pdf>.
- Zhenzhong Lan, Mingda Chen Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. A lite bert for self-supervised learning of language representations, 2020. URL <https://arxiv.org/pdf/1909.11942.pdf>.
- Junyi Li, Bin Wang, and Haiyan Ding. An albert model based maximum ensemble with different training sizes and depths for commonsense validation and explanation, 2020. URL https://www.researchgate.net/publication/355434928_Lijunyi_at_SemEval-2020_Task_4_An_ALBERT_Model_Based_Maximum_Ensemble_with_Different_Training_Sizes_and_Depths_for_Commonsense_Validation_and_Explanation.
- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space, 2013. URL <https://arxiv.org/pdf/1301.3781.pdf>.
- E. Turner. An introduction to transformers, 2024. URL <https://arxiv.org/pdf/2304.10557.pdf>.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2023. URL <https://arxiv.org/pdf/1706.03762.pdf>.
- Fuzhen Zhuang, Zhiyuan Qi, Keyu Duan, Dongbo Xi, Yongchun Zhu, Hengshu Zhu, Hui Xiong, and Qing He. A comprehensive survey on transfer learning, 2020. URL <https://arxiv.org/pdf/1911.02685.pdf>.